

Universitatea Tehnica din Cluj-Napoca
Departament Calculatoare

Programming Techniques in Java

Performance Tuning and Techniques

I. Salomie, C.Pop
2017

UTCN - Programming Techniques

1

Introduction

- Common perception: Java is slow ??!!
 - Many assume that if a program is not compiled it must be slow
 - Some of early Java versions (non-optimal coding, un-optimized JVMs, etc.) were slow
- Java overhead
 - JVM layer - abstracts Java from the physical machine
 - Java main advantages add certain performance overhead
 - Platform independence, memory management,
 - Exception checking,
 - Dynamic resource loading,
 - Security checks, etc.
 - OO and inheritance
 - Run time linking
 - Hierarchical method invocation
 - Small methods such as getters and setters - inline optimizations
- All first run programs (no matter implementation language) could seem slow
 - Not yet performance tuned
 - Bottlenecks not yet identified

UTCN - Programming Techniques

2

System limitations

- Resources that may limit performance for all applications
 - CPU speed and availability
 - System memory
 - Disk (and network) input/output (I/O)
- Question - (when tuning)
 - which of these resources is causing the application to run too slowly

System limitations

CPU speed and availability

- Look for possible improvements in the code
 - bottlenecks,
 - inefficient algorithms,
 - too many short-lived objects
 - Object creation and garbage collection are CPU-intensive operations
 - etc.

System memory limits

- Memory blocked with
 - Too many objects, or
 - A few large objects,
 - Too many allocated large arrays
 - Frequently used in buffered applications
- Check application design
 - May need to be reexamined to reduce its running memory footprint

Tuning strategy

Overview

- Identify the main bottlenecks
 - look for a few bottlenecks possibilities
- Choose the quickest and easiest one to fix and address it
- Repeat for other identified bottlenecks

Details

Loop the following sequence of actions:

- Measure the performance
 - use profilers and benchmark suites
 - instrumenting code
- Identify bottleneck locations
- Think of a hypothesis for the cause of the bottleneck
 - Consider any factors that may refute your hypothesis.
- Create a test to isolate the factor identified by the hypothesis.
 - Test the hypothesis.
- Alter the application to reduce the bottleneck.
- Test
 - the alteration improves performance
 - measure the improvement
 - do regression-testing the affected code

Tuning strategy

Profiler [Wikipedia]

- Profiling ("program profiling", "software profiling") is a form of [dynamic program analysis](#) that measures, for example, the space (memory) or time [complexity of a program](#), the [usage of particular instructions](#), or frequency and duration of function calls
- Profilers are used in the [performance engineering](#) process
- The most common use of profiling information is to aid program [optimization](#)
- Profiling is achieved by instrumenting either the program [source code](#) or its binary executable form using a tool called a *profiler* (or *code profiler*).
- A number of different techniques may be used by profilers, such as event-based, statistical, instrumented, and simulation methods

Tuning strategy

Profiling techniques

- Flat profilers (determine only the average time calls)
- Call-graph profilers (call times and frequencies of function calls, call chains, etc.)
- Event-based profilers – provide hooks to trap events, class load, unload, thread enter, leave, etc.
 - For Java: [JVMTI](#) (JVM Tools Interface) API, formerly JVMPPI (JVM Profiling Interface)
- Statistical based profilers
- Instrumented profilers
- Simulation based profilers

Tuning strategy

Recommended Java Profilers

- Jprofiler
<https://www.ej-technologies.com/products/jprofiler/overview.html>
- VisualVM
<https://visualvm.github.io/>
- Java Mission Control
<http://www.oracle.com/technetwork/java/javaseproducts/mission-control/java-mission-control-1998576.html>
- YourKit
<https://www.yourkit.com/>

Tuning strategy

Program Instrumentation

- Program instrumentation - modifies a program to analyze itself
 - The inserted code outputs the analysis data
 - Usually slows down the instrumented program
- Types
 - Manual inserted statements
 - Tools for automatic adding of the instrumented code
 - Compiler assisted adding of instrumented code, for example
 - `gcc -pg`
 - Runtime instrumentation- before the execution, the code is instrumented. Program execution is fully supervised and controlled by such a tool (ex. Pin, Valgrind, DynamoRIO);
 - Runtime code injection (more lightweight than Runtime instrumentation) – code modified at runtime. Inserts jumps to helper functions (ex. DynInst) ⁹

UTCN - Programming Techniques

Go faster

Threading to appear quicker

- Use multithreading to make an application responsive
- Use threads to ensure that any particular service is available and unblocked when needed
 - Difficult to program correctly and manage
 - Complex task - handling inter-thread communication

UTCN - Programming Techniques

10

Go faster

Streaming to appear quicker

- The general case
 - A long activity that can provide results in a stream is identified
 - The results can be accessed a few at a time
- The case for distributed applications
 - Sending all data vs. streaming data while it is processed
 - Example:
 - Web browsers - display the initial page screen as soon as it is available, without waiting for the whole page to be downloaded

UTCN - Programming Techniques

11

Go faster

Caching to Appear Quicker

- Moving data is expensive
 - If the same data is used multiple times => keep a copy
 - Example: browser downloads
 - Disk access example:
 - from lowest-level hardware cache
 - to OS read / write cache
- Cached file systems
 - File reading / writing classes that provide buffered I/O
 - Cost of reading a byte is the same with reading a page

UTCN - Programming Techniques

12

Performance and Metrics

- Applications must meet user expectations
 - Functional requirements
 - Non-functional requirements
 - ... performance among them
- Agree on performance targets
 - Agree on performance metrics
 - Identify target response times for as many subsystems as possible
- Without clear performance objectives tuning will never be completed
- System-wide throughput
 - number of transactions per minute for the system as a whole
 - response times on a saturated network
- The maximum number of resources the application should support
 - users, data, files, file sizes, objects

UTCN - Programming Techniques

13

Setting benchmarks

- Benchmark (Wikipedia)
 - The result of running a computer program, or a set of programs, in order to assess the relative performance of an object by running a number of standard tests and trials against it
- Setting benchmarks - precise specification stating
 - High level - what functionality needs to run in what amount of time
 - Low level - what part of the code needs to run in what amount of time
 - Specify in terms of user actions
 - the time elapsed from user pressing the button XXX until a result is displayed on the screen
- Important questions
 - How fast ?
 - In which parts ?
 - For how much effort ?

UTCN - Programming Techniques

14

Setting benchmarks (cont.)

- Target times specification for each benchmark
 - best times, accepted times, etc.
 - Example
 - Function X takes to execute - from pressing the button until a response is displayed
 - 2 sec in 70% of the cases
 - 2-3 sec in 15% of the cases
 - 3-5 sec in 15% of the cases
 - Anything that can vary should be controlled and reproducible

Benchmark Harness

- Tools for testing applications
- Benchmark harness - different complexities
 - from a simple class that sets some values and starts main()
 - logging and timestamp
 - to GUI-run benchmark harness
- Examples
 - SPECjbb2013 from
 - <http://www.spec.org/jbb2013/>
 - <http://www.spec.org/osg/jbb2005/>

Benchmark Harness

Requirements

- Correctly reproduce user activity and data input and output
 - Reproduce and simulate all user input, including GUI input
- Support different application configurations
- Control any randomized inputs
- Random sequences used in tests should be reproducible
- Support for logging statistics
- Allow testing the system across all scales of intended use
 - up to the maximum # of users, simulate users
 - objects
 - throughputs, etc
- Each run needs to be under identical conditions (as much as possible)

UTCN - Programming Techniques

17

Benchmark Harness

How to measure

- The benchmarks should be run multiple times
 - the full list of results should be stored and processed
 - i.e. not just the average and deviation or the ranged percentages
 - Assure similar conditions for benchmark runs
- Each code change should be driven by
 - Identifying exactly which bottleneck is to be improved
 - How much speedup is expected
- After each code change (or set of changes)
 - Run benchmarks to precisely identify improvements (or degradations) in the performance across all functions
- Code Optimizing
 - Can introduce new bugs, so the application should be tested during the optimization phase.
 - Validated only after the application using that optimization's code path has passed quality assessment.
 - Should be completely documented

UTCN - Programming Techniques

18

Benchmark Harness

What to measure

- Main measurement reports
 - time
 - throughput
- Main way to measure time
 - **System.currentTimeMillis()**
 - two calls to determine the elapsed time
 - works well for not short periods of time
 - Call costs
 - takes up to 1/2 millisecond to execute
 - for a period of 100 milliseconds the call cost is 1%
 - Other method
 - **System.nanoTime()**
- Always small variations between test runs
 - use averages to measure differences and
 - consider whether those differences are relevant by calculating the variance in the results.
- CPU time
 - time allocated on the CPU for a particular procedure
- The # of runnable processes waiting for CPU
 - this gives you an idea of CPU contention
- Memory sizes
- Disk throughput
- Disk scanning times
- Network traffic, throughput, and latency
- Transaction rates
- Other system values
- Note. Measuring these values needs:
 - system knowledge - no Java mechanisms for direct measuring of these values
 - application-specific knowledge (what is a transaction for your application?)

UTCN - Programming Techniques

19

Techniques for Performance

Loops

- Move out code from loops
- Code that does not need to be executed on every pass
- Assignments, accesses, tests, and method calls that need to run only once
- Method call has costs
 - Analyze task - moving method calls out of the loop (even if this requires rewriting)

UTCN - Programming Techniques

20

Techniques for Performance

Loops

- Termination test and method calls
 - Avoid method call in a loop termination test
 - Significant overhead

- Example

```
for(int i = 0; i < collection.size( ); i++)
```

- Factor out the method call

```
int max = collection.size( );
```

```
for(int i = 0; i < max; i++)
```

UTCN - Programming Techniques

21

Techniques for Performance

Loops

- Loops, temporary variables and arrays
 - VM performs bounds-checking for array-element access
 - Array access (and assignment)
 - More overhead than temporary variable access
 - Inefficiency - repeated access in each iteration
 - Do array access once - and assigned to a temporary outside the loop

- Example

```
for(int i = 0; i < repNo; i++)
  countArr[0] += 10;
```

- Optimized loop

```
count = countArr[0];
for(int i = 0; i < repNo; i++) count += 10;
countArr[0] = count;
```

Note. Twice as fast

UTCN - Programming Techniques

22

Techniques for Performance

Loops

- Use int for Index Variables
 - int is the fastest data type
 - The VM is optimized to use ints
 - Operations on bytes, shorts, and chars - implicit casts to/from ints

- Example

Instead of

```
for(long i = 0; i < repNo; i++)
```

```
for(char i = 0; i < repNo; i++)
```

use int data types

```
for(int i = 0; i < repNo; i++)
```

UTCN - Programming Techniques

23

Techniques for Performance

Loops

- Loops vs. System.arraycopy()
 - Copy arrays
 - Loop
 - System.arraycopy()
 - System.arraycopy() is faster
 - when VM has a JIT - loop could be little faster
- Conclusion
 - always use System.arraycopy()
 - even if a VM JIT is used

UTCN - Programming Techniques

24

Techniques for Performance

Loops

- Comparison to 0 is faster than comparisons to most other numbers.
 - The VM has optimizations for comparisons to the integers
 - -1, 0, 1, 2, 3, 4, and 5.
 - Rewrite loops to make the test a comparison against 0
 - Note. The latest VMs try to optimize the standard loop expression, so rewriting the loop may not produce faster code
- Reverse the iteration order of the loop
 - from counting up (0 to max) to counting down (max to 0).
 - Instead of


```
for(int i = 0; i < repNo; i++)
```
 - Could be coded (for better performance):


```
for(int i = repNo-1; i >= 0; i--)
for(int i = repNo; --i >= 0 ; )
```

UTCN - Programming Techniques

25

Techniques for Performance

Exceptions

- Exception generation know-how
 - important for good design and
 - performance achievements
 - an exception should be thrown only in exceptional situations
 - Example
- try-catch blocks costs
 - no extra time if no exception is thrown
 - some VMs may impose a slight penalty
 - significant overhead if an exception is thrown
 - due to stack snapshot creation
 - equivalent to the execution of several hundred of simple code lines

UTCN - Programming Techniques

26

Techniques for Performance Exceptions

- Implicitly generated exceptions
 - JVM generated exceptions
 - ClassCastException
 - ArrayIndexOutOfBoundsException
- Explicitly thrown exceptions
 - using throw statement
 - StackTrace generation overhead
- Reusing an existing exception object
 - reduce cost by not creating a new exception object
 - two orders of magnitude faster
 - avoids stack trace generation => reduce the overhead

UTCN - Programming Techniques

27

Techniques for Performance Exceptions

Reusable Exceptions - Example

```
public static Exception
    REUSE_EXCEPTION =
        new Exception( );
```

...

//Faster by exception reusing

```
try {
    ... throw REUSE_EXCEPTION;
}
catch (Exception e) {...}
```

//The following try-catch is 50..100

// times slower

```
try {
    ... throw new Exception( );
}
catch (Exception e) {...}
```

- Disadvantage
 - the stack trace is incorrect for the thrown exception object
 - it is the one that is generated when the exception is created
 - this could be important and may generate maintenance future problems
- Note
 - fillInStackTrace() should be invoked to get the stack filled with the current situation
 - this call generates the mentioned overhead!!

UTCN - Programming Techniques

28

Techniques for Performance

try/catch vs. instanceof

// pesimistic approach

```
public static boolean test1(Object obj) {
    if (obj instanceof Integer) {
        Integer i = (Integer) obj;
        return false;
    }
    else
        return true;
}
```

// optimistic approach (speculative cast)

```
public static boolean test2(Object obj) {
    try {
        Integer i = (Integer) obj;
        return false;
    }
    catch (Exception e) {return true;}
}
```

Conclusion

- test1() takes 1 TU (assumption)
- test2() with ClassCastException thrown
 - 100 TU
- Recommendation
 - use instanceof instead of speculative cast

Techniques for Performance

Method parameter checking

- Check method arguments for validity

- Ok during development / testing time
- Significant overhead for the released product
- Error checking could be removed for the released product

- Next is a technique for error parameter checks that can be optionally removed (through an extra compilation)

// Conditional Error Checking)

```
public class GLOBAL_CONSTANTS {
    public static final boolean ERR_CHECK_ON = true;
    ...
}
```

```
//code in methods of other classes
if (GLOBAL_CONSTANTS.ERR_CHECK_ON) {
    //error check code of some sort
    ...
}
```

- Allows turning ON/OFF error checking by recompiling
- A compiler feature – eliminates the overhead of the if statement for each check

Techniques for Performance

Assertions

- Forms
 - assert boolean_expression;**
 - if boolean expression evaluates to false => AssertionError is thrown.
 - assert boolean_expression : String_expression;**
 - allows message customization using a String parameter
 - Example


```
assert val >= 0 : "Parameter val must be non-negative, but was " + val
```
- assert statement - normal statement comparison
 - assert can be enabled / disabled at runtime
- Assertions enabled
 - Better quality code
 - Assists in diagnosing problems
- Assertion enable / disable
 - Assertion enable / disable granularity - class level
 - **-ea** and **-da** parameters of **java** command can be used to specify for each class if assertions are enabled or disabled

UTCN - Programming Techniques

31

Techniques for Performance

Assertions

How assertion works

- At compile time
 - When an assertion is found in a class
 - A new field is added to the class - and is left unassigned (internal legal operation)
 - \$assertionsDisabled
 - How assertion is compiled


```
if (! $assertionsDisabled)
  if (!boolean_expression)
    throw new AssertionError(String_expression);
```
- At class loading time
 - class loader determines assertion status of the class using the specified assertion rules and sets the value of \$assertionsDisabled

UTCN - Programming Techniques

32

Techniques for Performance Assertions

Conclusion

- Assertions are disabled
 - Without any further optimization every assert costs a minimum one runtime test
 - Significant overhead for short methods
 - Setters, getters and other short frequently called methods - think of avoiding
- Assertions are enabled
 - Any assertion takes at least as long to run as its `boolean_expression` evaluation takes
 - Code running with assertions enabled will be slower than code running with assertions disabled - even if only a few percent slower

Techniques for Performance Assertions

Assertion vs. explicitly checks

- Facts
 - Assertions can be turned off
 - Explicit checks cannot be turned off
- Conclusion
 - Consider changing all explicit checks for incorrect parameters and state in your code to use assertions instead of explicitly using **if...throw** statements
 - Example
 - `IllegalArgumentException`s often test for documented incorrect conditions
 - These tests could be changed to assertions
- Ultimate decision
 - The test should always be present
 - Don't make it an assertion
 - The test is optional and provides extra robustness, especially during development and testing
 - Use assertion

Techniques for Performance Assertions

Examples

```
public int noCheckMethod(int iVal) {
    // ... method body
    return ...
}

public int explicitlyCheckMethod(int iVal) {
    if(iVal < 0)
        throw new IllegalArgumentException("parameter val should be positive, but is " +
            val);
    // ... method body
    return ...
}

public int usingAssertMethod(int iVal) {
    assert (val >= 0) : "parameter val should be positive, but is " + val;
    // ... method body
    return ...
}
```

UTCN - Programming Techniques 35

Techniques for Performance Casts

- Resolve time
 - Some casts can be resolved at compile time
 - They are eliminated by the compiler
 - Casts not resolvable at compile time must be executed at runtime
- Casts cost dependences
 - Depends on the depth of the hierarchy
 - The further back in the hierarchy the longer the cast takes to execute
 - Depends whether the casting type is an interface or a class
 - Interfaces are generally more expensive to use in casting
- Important note
 - Never change the design of the application for minor performance gains

Techniques for Performance Casts

Good practices

- Avoid casts whenever possible
 - Use type-specific collection classes instead of generic collection classes
 - Instead using a standard List to store a list of Strings
 - Better performance with a StringList class
 - Type the variable as precisely as possible
- Variable needing casting several times
 - Cast once and save the object into a temporary variable of the cast type
 - Use the temporary variable instead of repeatedly casting

UTCN - Programming Techniques

37

Techniques for Performance Casts

Bad practices

- Avoid the following type of code


```
if (obj instanceof Something)
    return ((Something)obj).x + ((Something)obj).y + ((Something)obj).z;
...

• Use a temporary variable
    – more readable code
    if (obj instanceof Something) {
        Something something = (Something) obj;
        return something.x + something.y + something.z;
    }
```

... UTCN - Programming Techniques

38

Techniques for Performance Variables

- Local (temporary) and method-argument variables
 - The fastest variables to access and update
 - Local variables
 - Remain on the stack, so they can be manipulated directly
 - There are special bytecodes for accessing the first four local variables and parameters on a method stack
- Static and instance variables
 - Stored and manipulated in the heap memory
 - Instance and static variables vs. local variables and method arguments
 - Up to an order of magnitude slower to operate on

UTCN - Programming Techniques

39

Techniques for Performance Variables

Temporary objects

- Created from the heap
- Object reference itself is allocated on the stack
- Operations on any object are slower than on any of the primitive data types for temporary variables
- Temporary primitive variables vs. temporary objects
 - As soon as variables are discarded at the end of a method call, the memory from the stack can immediately be reused for other temporaries.
 - Any temporary objects remain in the heap until garbage collection reallocates the space
- **Conclusion:** Temporary variables using primitive (non object) data types are better for performance

UTCN - Programming Techniques

40

Techniques for Performance Variables

Good practice

- When many manipulations on instance or static variable
 - Execute them on a temporary variable
 - Reassign to the instance variable at the end
 - This is true for instance variables that hold arrays as well
- Where possible
 - Access public instance variables rather than use methods (getters and setters)
 - This breaks encapsulation (bad design)
 - Java SDK uses this techniques in some of the cases

UTCN - Programming Techniques

41

Techniques for Performance Variables

Good practice

- Array-element access - typically two to three times as expensive as accessing non-array elements
 - Due to range checking and null pointer checking (for the array itself) done by the VM
 - The VM JIT compiler manages to eliminate almost all the overhead in the case of large arrays
- Arrays overhead, due to the range checking in Java
 - If an array element is manipulating many times => assign it to a temporary variable for the duration

```
for(int i = 0; i < repNo; i++)
    countArr[0] += i;
```

Replace with:

```
int count = countArr[0];
for(int i = 0; i < repNo; i++) count += i;
countArr[0] = count;
```

UTCN - Programming Techniques

42

Techniques for Performance Variables

Good practice

- ints
 - The fastest variable type to operate on
- shorts, bytes, and chars
 - All widened to ints for almost any type of arithmetic operation
 - Cast back is required if you want to end up with the data type you started with
 - Example
 - Adding two bytes produces an int and requires a cast to get back a byte
- longs and doubles
 - Can take longer to access and update than other variables
 - Due to ... they are twice the basic storage length for Java (which is four bytes)
- Floating-point arithmetic
 - The worst in terms of performance